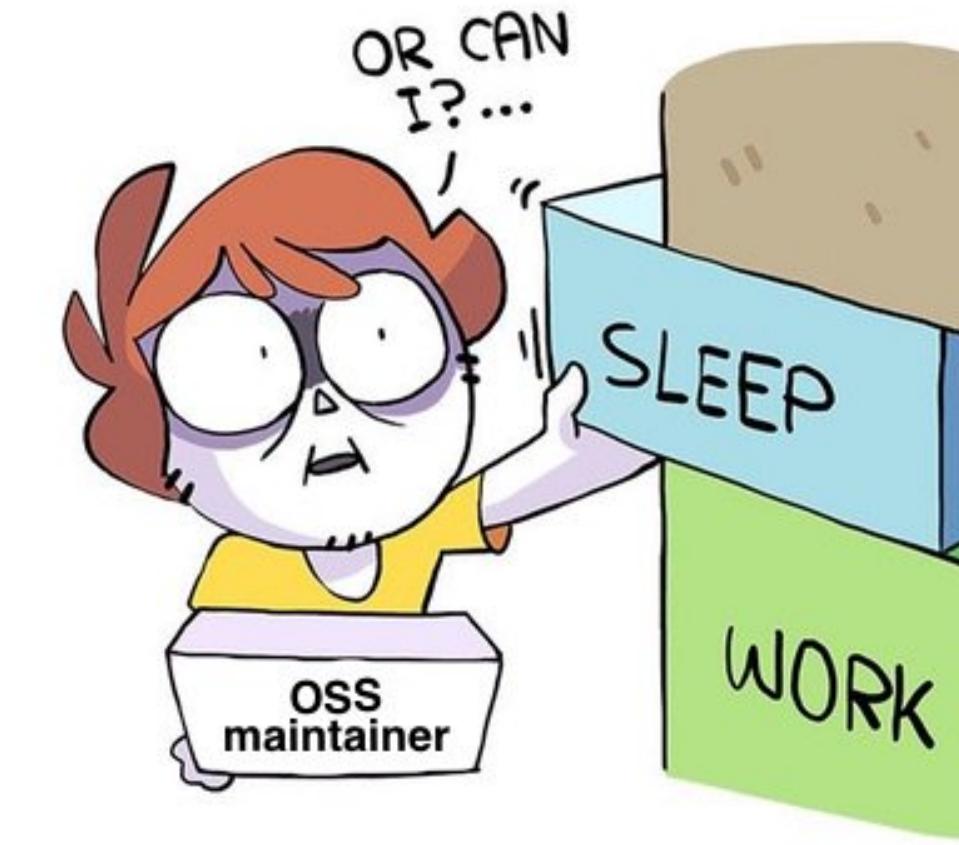
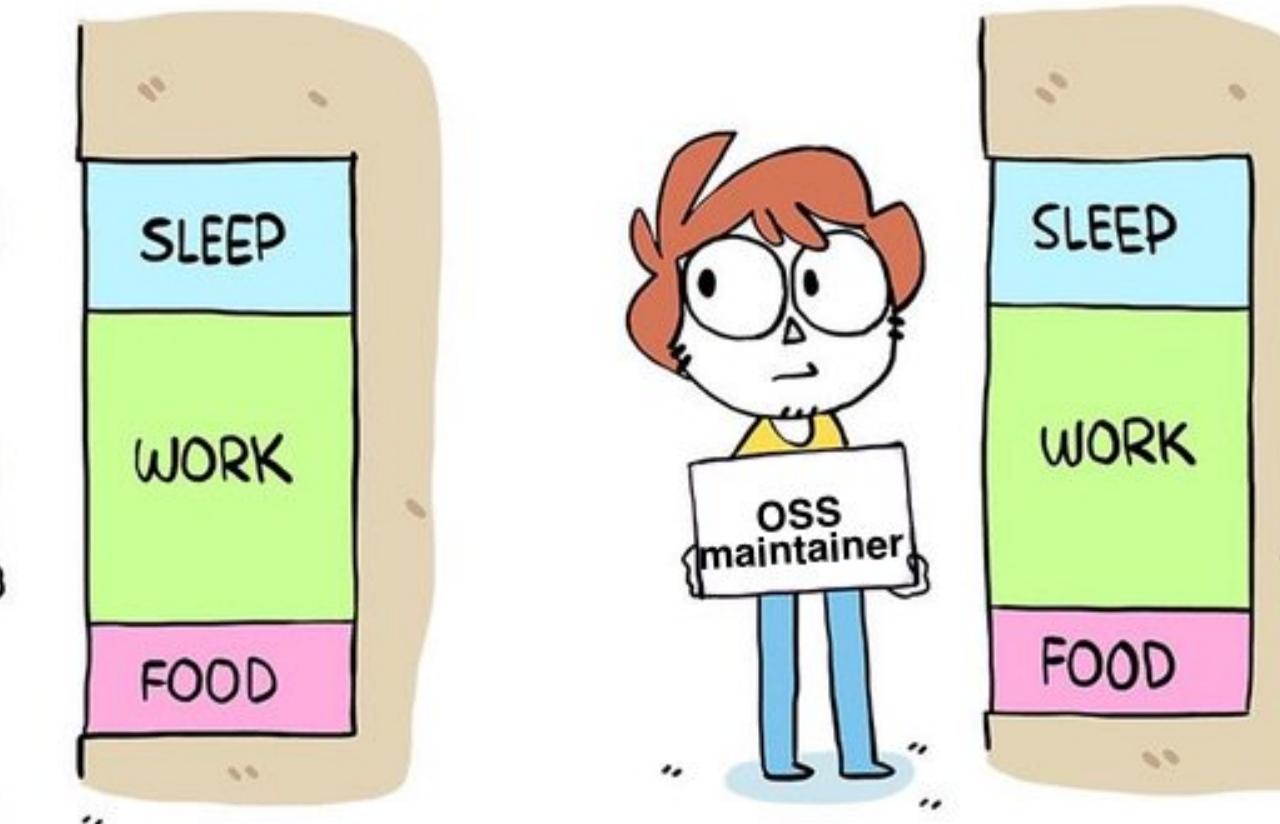


MEMORY FRAGMENTATION AND BLOAT IN RUBY

NATE BERKOPEC

SPEEDSHOP.CO



speedshop

Speedshop is a Ruby on Rails performance consultancy that optimizes the full stack - frontend, backend and environment - to generate revenue and cut scaling costs for businesses on Rails. It's not uncommon for Speedshop clients to halve their server costs and improve perceived load times by 300%. Fast sites are profitable sites. Speed is a feature.

Products and Services

- [The Complete Guide to Rails Performance](#), an in-depth 370 page book and 17 hours of HD screencasts and interviews with top experts, including DHH himself.
- [Speedshop Blog](#), the number one Ruby on Rails performance blog on the 'net.

Who is Speedshop?



Who would win

popular web framework



512MB of RAM



MEMORY CONSTRAINED ENVIRONMENTS

- * EMBEDDED**
- * 512 MB DYNOS**
- * 50% (?) OF RAILS APPS ARE
MEMORY-CONSTRAINED**

**MEMORY BEHAVIOR IS
DIFFICULT
FOR RUBYISTS TO
UNDERSTAND**

LAYERS OF ABSTRACTION

1. Your Ruby code itself
2. Ruby runtime
3. Allocator
4. MMU (virtual/real memory)
5. Actual memory location (swap/RAM)

**WE SHOULD PREVENT
MEMORY OVERUSE
OR MAKE IT
EASIER TO
DEBUG AND FIX**

BLONDA



PATTERNS

**SUDDEN SPIKES
IN MEMORY, SLOW/NO FALLOFF**

Latest

▲ 144.1 % 1,474.4 MB

Max

▲ 167.3 % 1,713.4 MB



**MEMORY USAGE THAT'S
NECESSARY, BUT
EXCESSIVE**

**PRIMARILY A FUNCTION
OF THE
RUNTIME AND YOUR
USER CODE**

**CONTRIBUTES TO RAILS
PROCESSES THAT NEED**

~1GB EACH

User.all.each



API::V1::ADMIN::EXPORTSCONTROLLER

```
def index  
  @every_single_user.to_csv  
end
```

“Note that, in general, “freeing” memory does not actually return it to the operating system for other applications to use.

However, if the top chunk in a heap - the portion adjacent to unmapped memory - becomes large enough, some of that memory may be unmapped and returned to the operating system.”

» [MallocInternals](#)

“Closing as NOTABUG based on comment 3. malloc/free indeed cannot give address space back to the kernel when main arena is discontinuous and that is expected behaviour.”

» Malloc bug 15321

SOLUTION[?]:

PAIN

IN INTERFACES



MINITEST: PAINFUL STUBS

```
Time.stub :now, Time.at(0) do
  assert objUnderTest.stale?
end
```

RAILS "STRICT MODE"

```
SomeModel.all.each # *always* batches  
SomeModel.select(:some_attribute) # *always* select
```

ENUMERABLE

```
raise CollectionTooLargeError # dev only, WARN in prod  
some_huge_collection.map # raise
```

**BUT THIS CAN BE
ABUSED:
WE DON'T WANT TO
"THINK LIKE A
COMPUTER!"**

“No user-level memory management
Ruby has automatic memory management.”

» Ruby manpage

**SOLUTION[?]:
JUST ALLOCATE LESS
SINATRA VS RAILS**

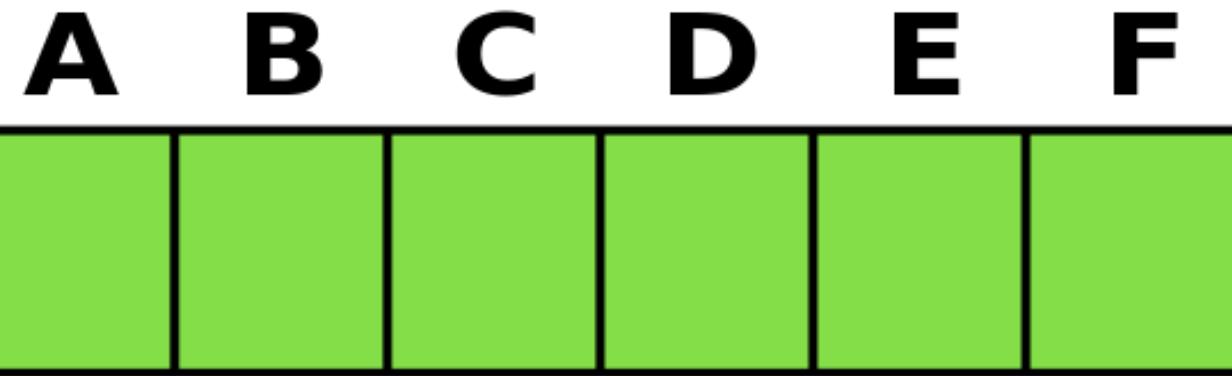
**SOLUTION?
AGGRESSIVE
DE-ALLOCATION**

FRAGMENTATION

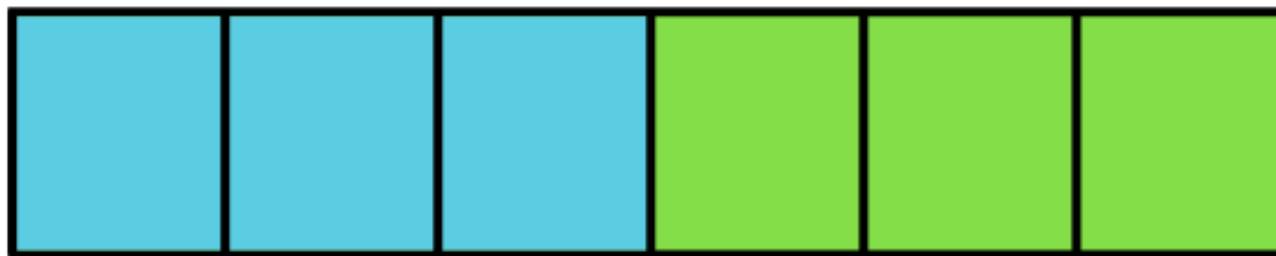
PATTERNS

**SLOW, ENDLESS
"LOGARITHMIC" GROWTH IN
LONG-RUNNING RUBY
PROGRAMS**

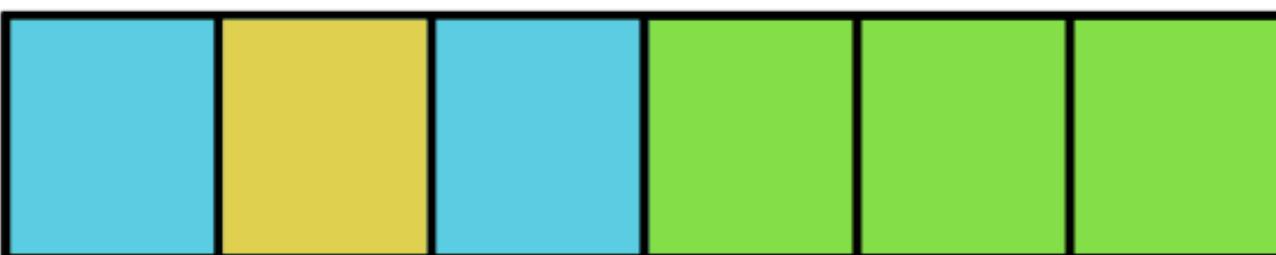
External fragmentation



Total free memory available for allocation.



Dynamically allocated
three blocks of memory (**A, B, C**).



Out of these three continuous blocks of
allocated memory, consider that the middle
block B is released. It is not possible to use
the freed block B, if the memory to be allocated
is larger than the size of block B.

SWISS

CHEESE

**WEB REQUEST
LOCATION
PATTERN**

PRIMARILY

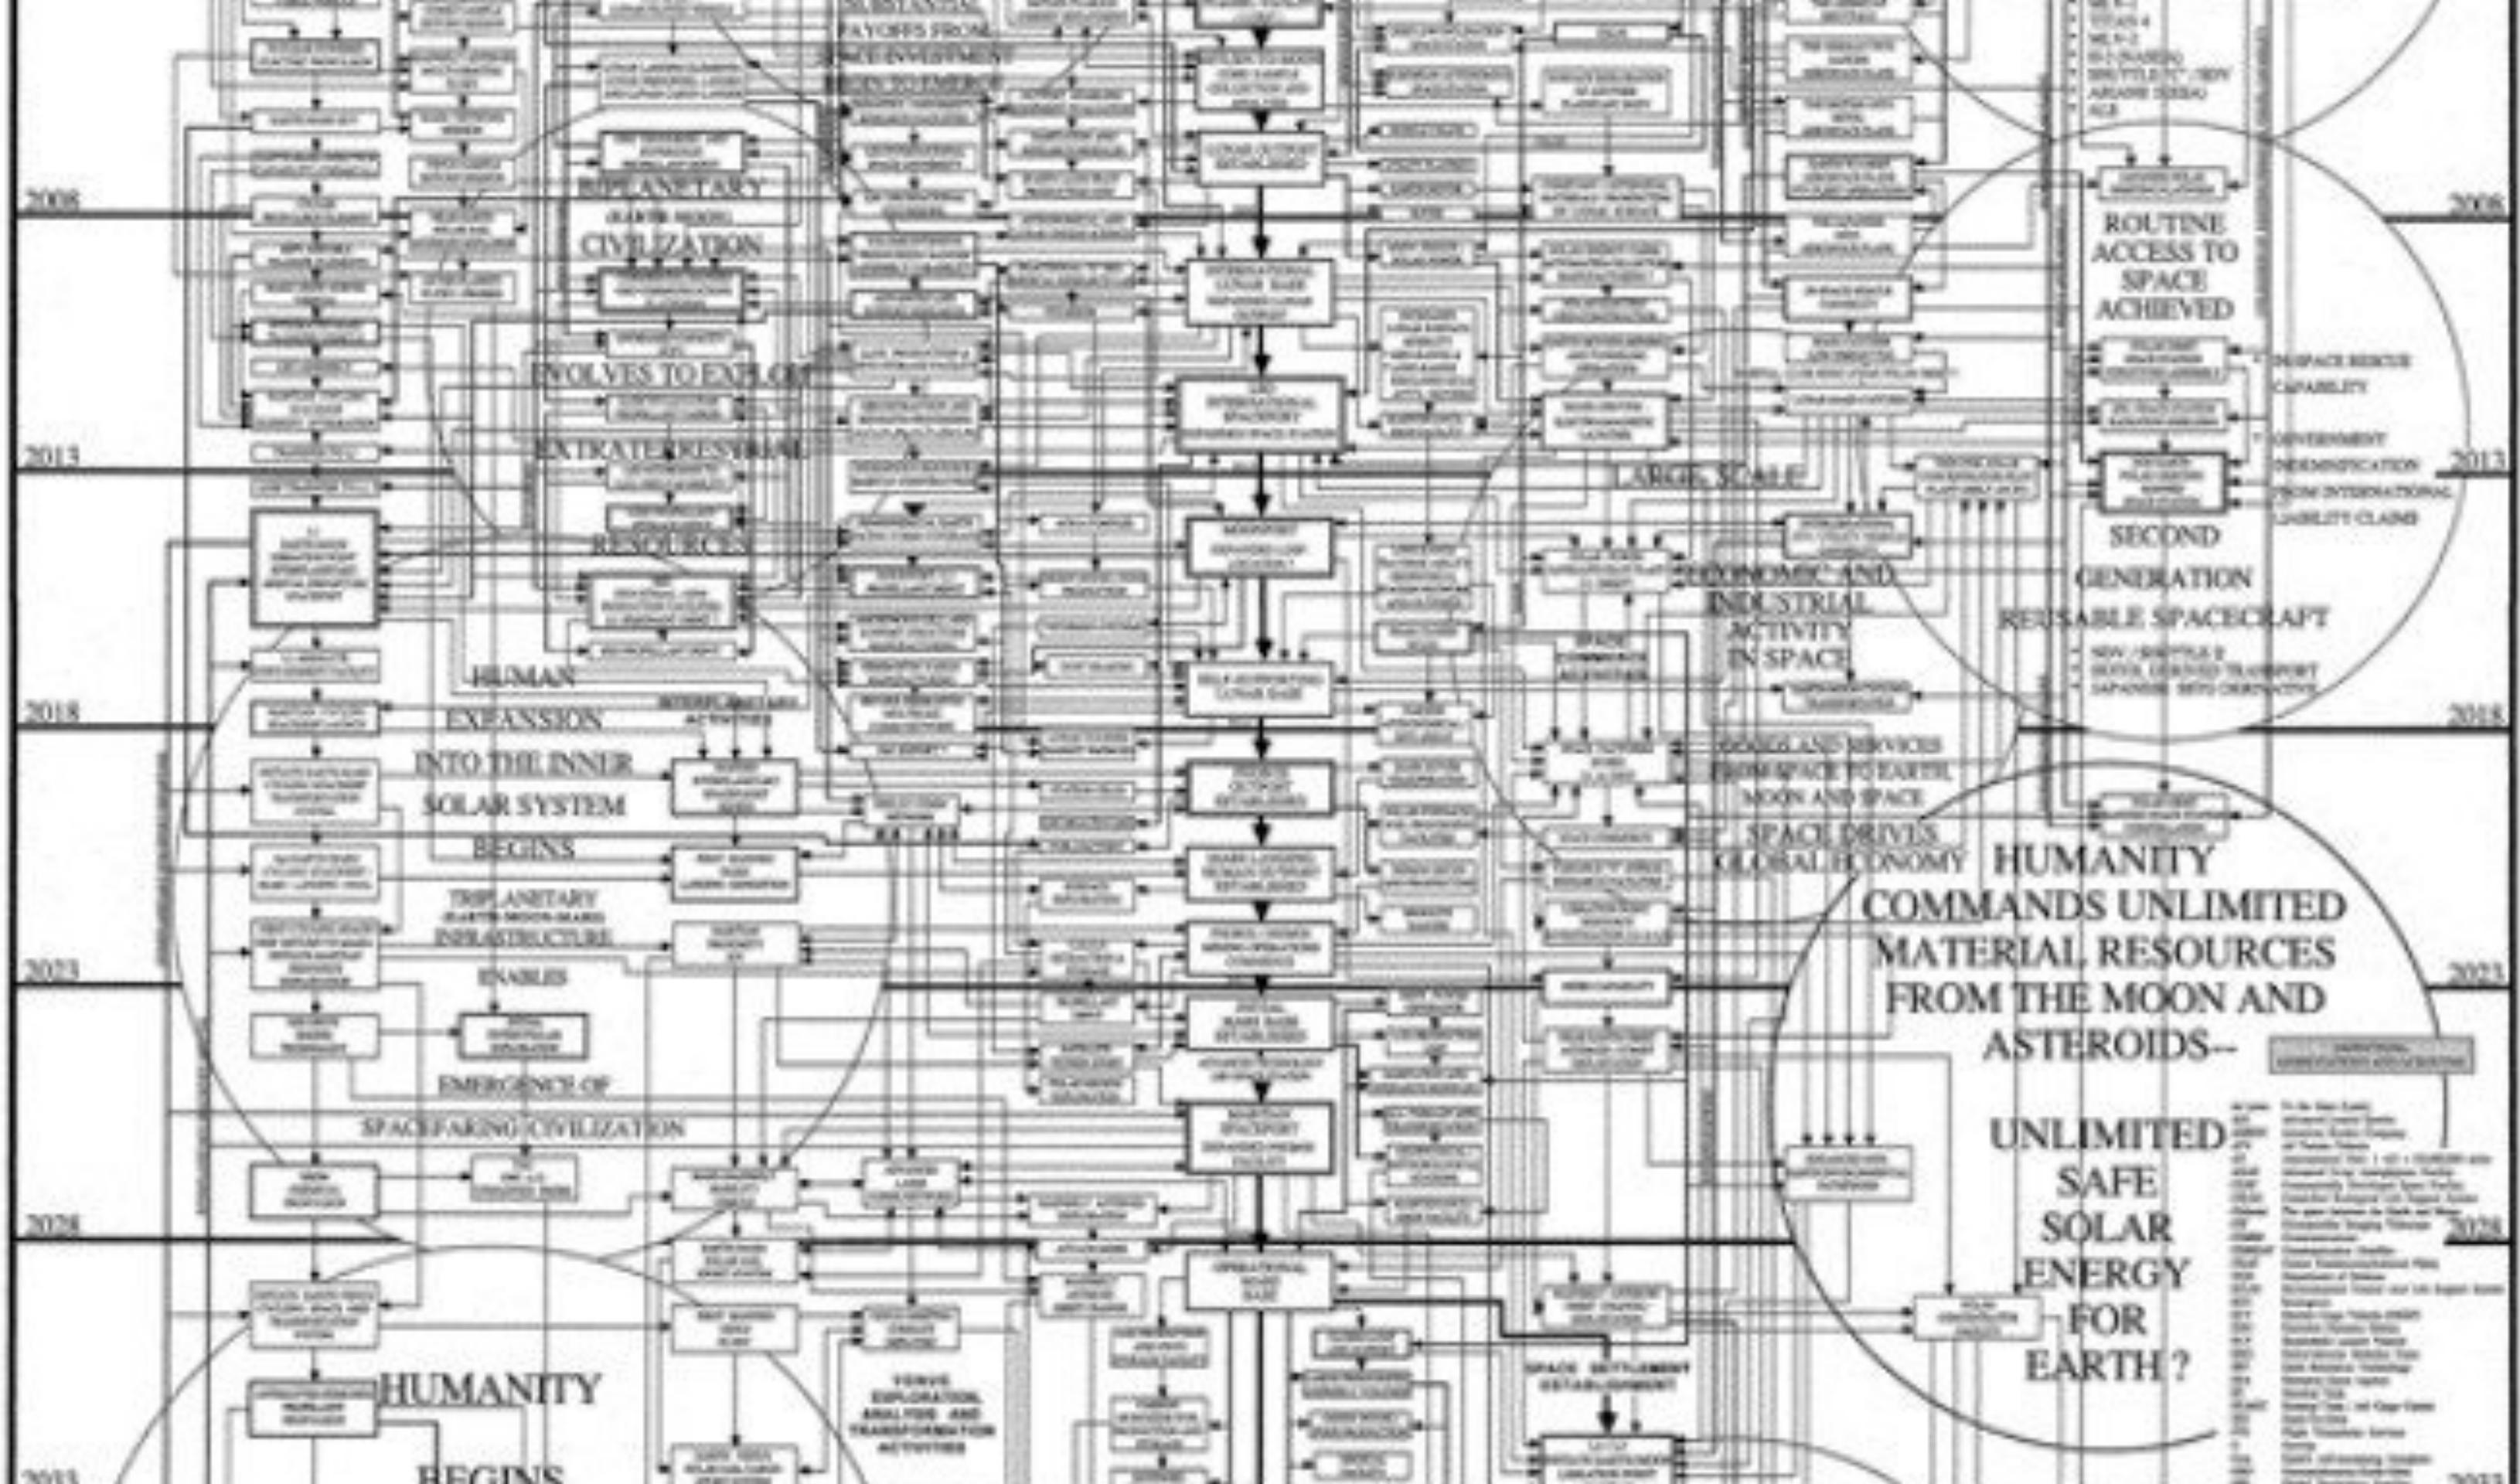
RUNTIME

AND ALLOCATOR

**OFTEN
MISTAKEN
FOR LEAKS
(LINEAR V LOG
GROWTH)**

GO START

**RUBY GC
INTERNALS IN
60 SECONDS**



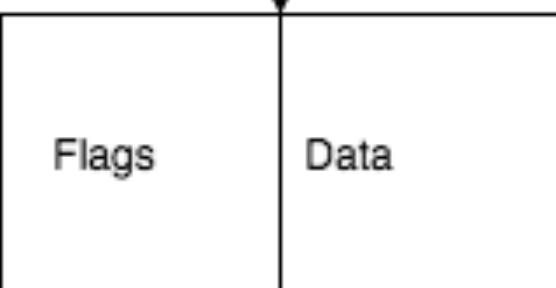
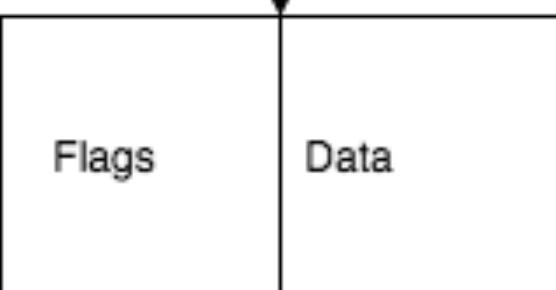
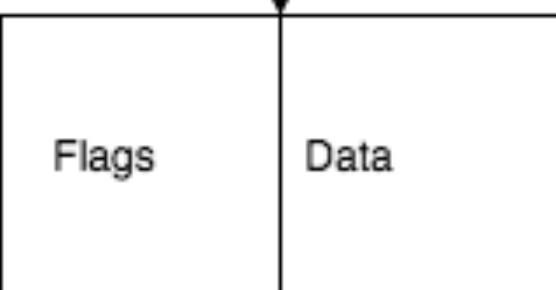
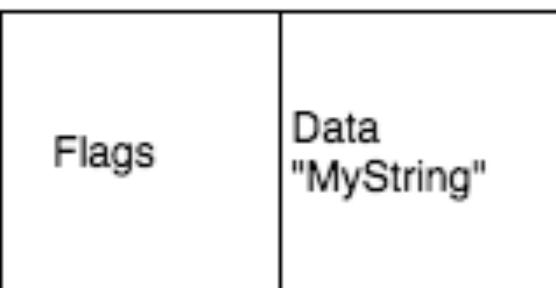
**EVERY OBJECT IS GIVEN
A 40-BYTE**

RVA VALUE

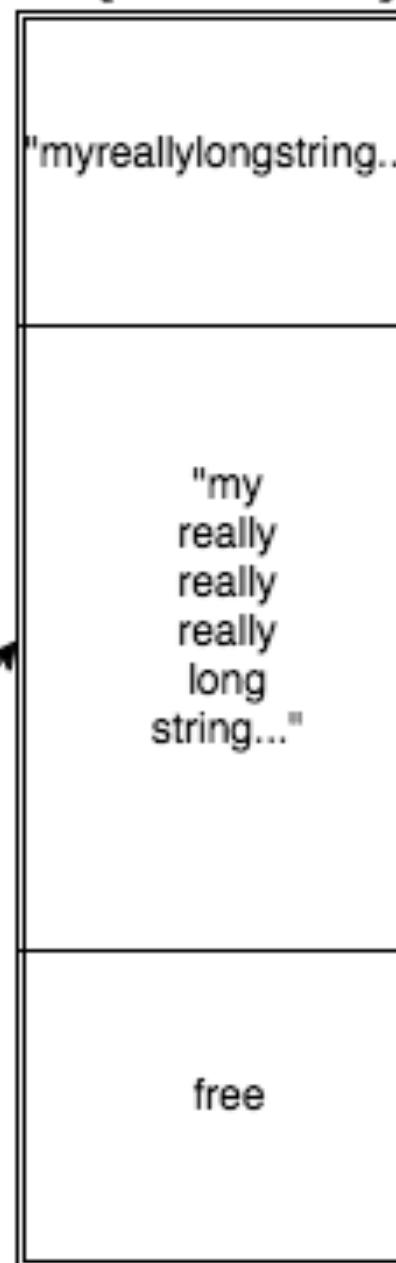
**ORGANIZED INTO
408-SLOT
"HEAP PAGES"**

**SOME (MOST?) OBJECTS
CONTAIN A POINTER TO
ALLOCATED MEMORY
WHERE ACTUAL OBJECT
DATA IS STORED**

RVALUE list (ObjectSpace)



C Heap (malloc)



```
5.times { GC.start }

GC.stat[:heap_live_slots] # 24508
GC.stat[:heap_eden_pages] # 83
GC::INTERNAL_CONSTANTS[:HEAP_PAGE_OBJ_LIMIT] # 408

# live_slots / (eden_pages * slots_per_page)
# 24508 / (83 * 408) = 72.3%
```

Page 1

Page 2

Page 3

Page 1

Page 3

```
GC.stat[:heap_sorted_length] # 83
```

```
# sorted_length / eden_pages
```

```
# 1
```

ALLOCATOR

PROBLEM:

PER-THREAD ARENAS IN GLIBC

"PUMA USES MORE
MEMORY THAN
UNICORN"

SIDEKIQIS
LEAKING

**PER-THREAD READ
MEMORY ARENAS
REDUCE CONTENTION
FOR THE MAIN ARENA**

“As pressure from thread collisions increases, additional arenas are created via mmap to relieve the pressure. The number of arenas is capped at eight times the number of CPUs in the system [...], but the trade-off is that there will be less fragmentation.”

» [MallocInternals](#)

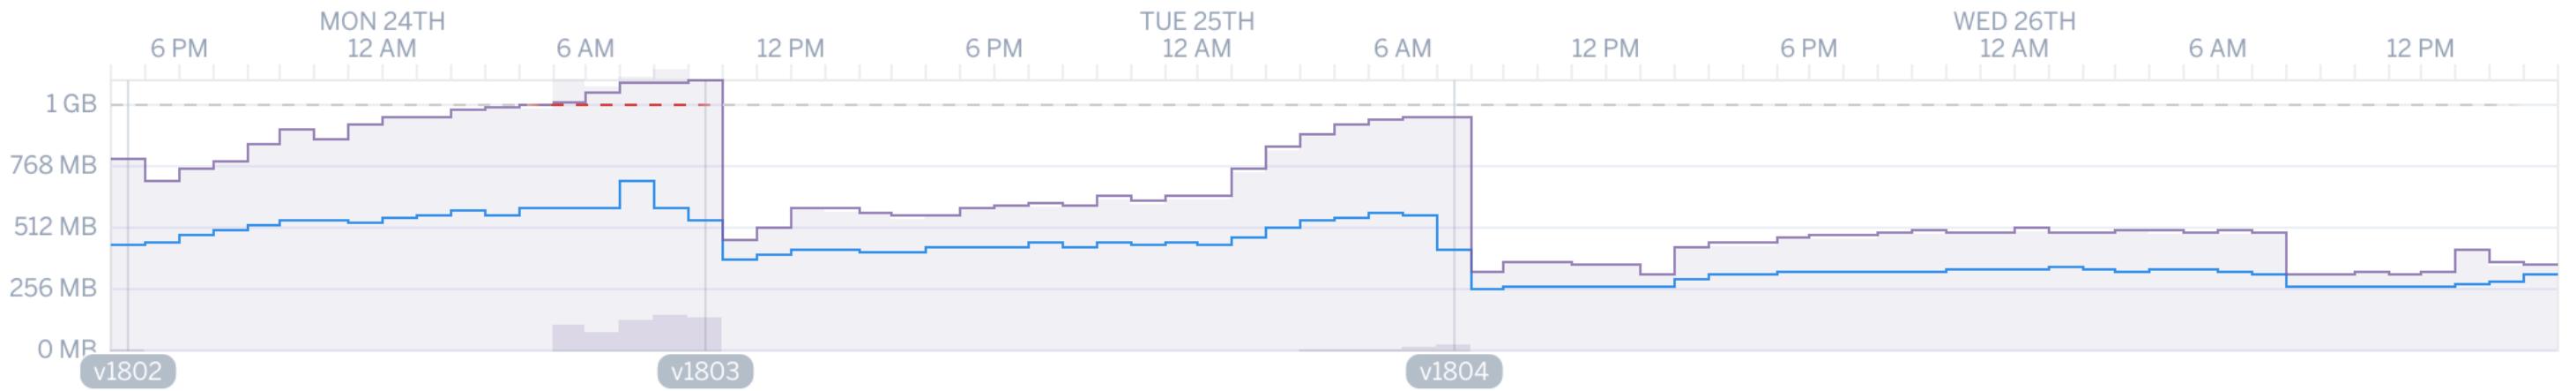
MALLOC_arena_max

Memory Usage i

Latest
35.1% 363.4 MB

Max
▲ 115.2% 1,179.8 MB

Average
40.2% 411.6 MB



- AVG. TOTAL
-
- MAX TOTAL
-
- MAX RSS
-
- MAX SWAP
-
- MEM QUOTA
-

**MAJOR SOURCE OF
RUBY
FRAGMENTATION?**

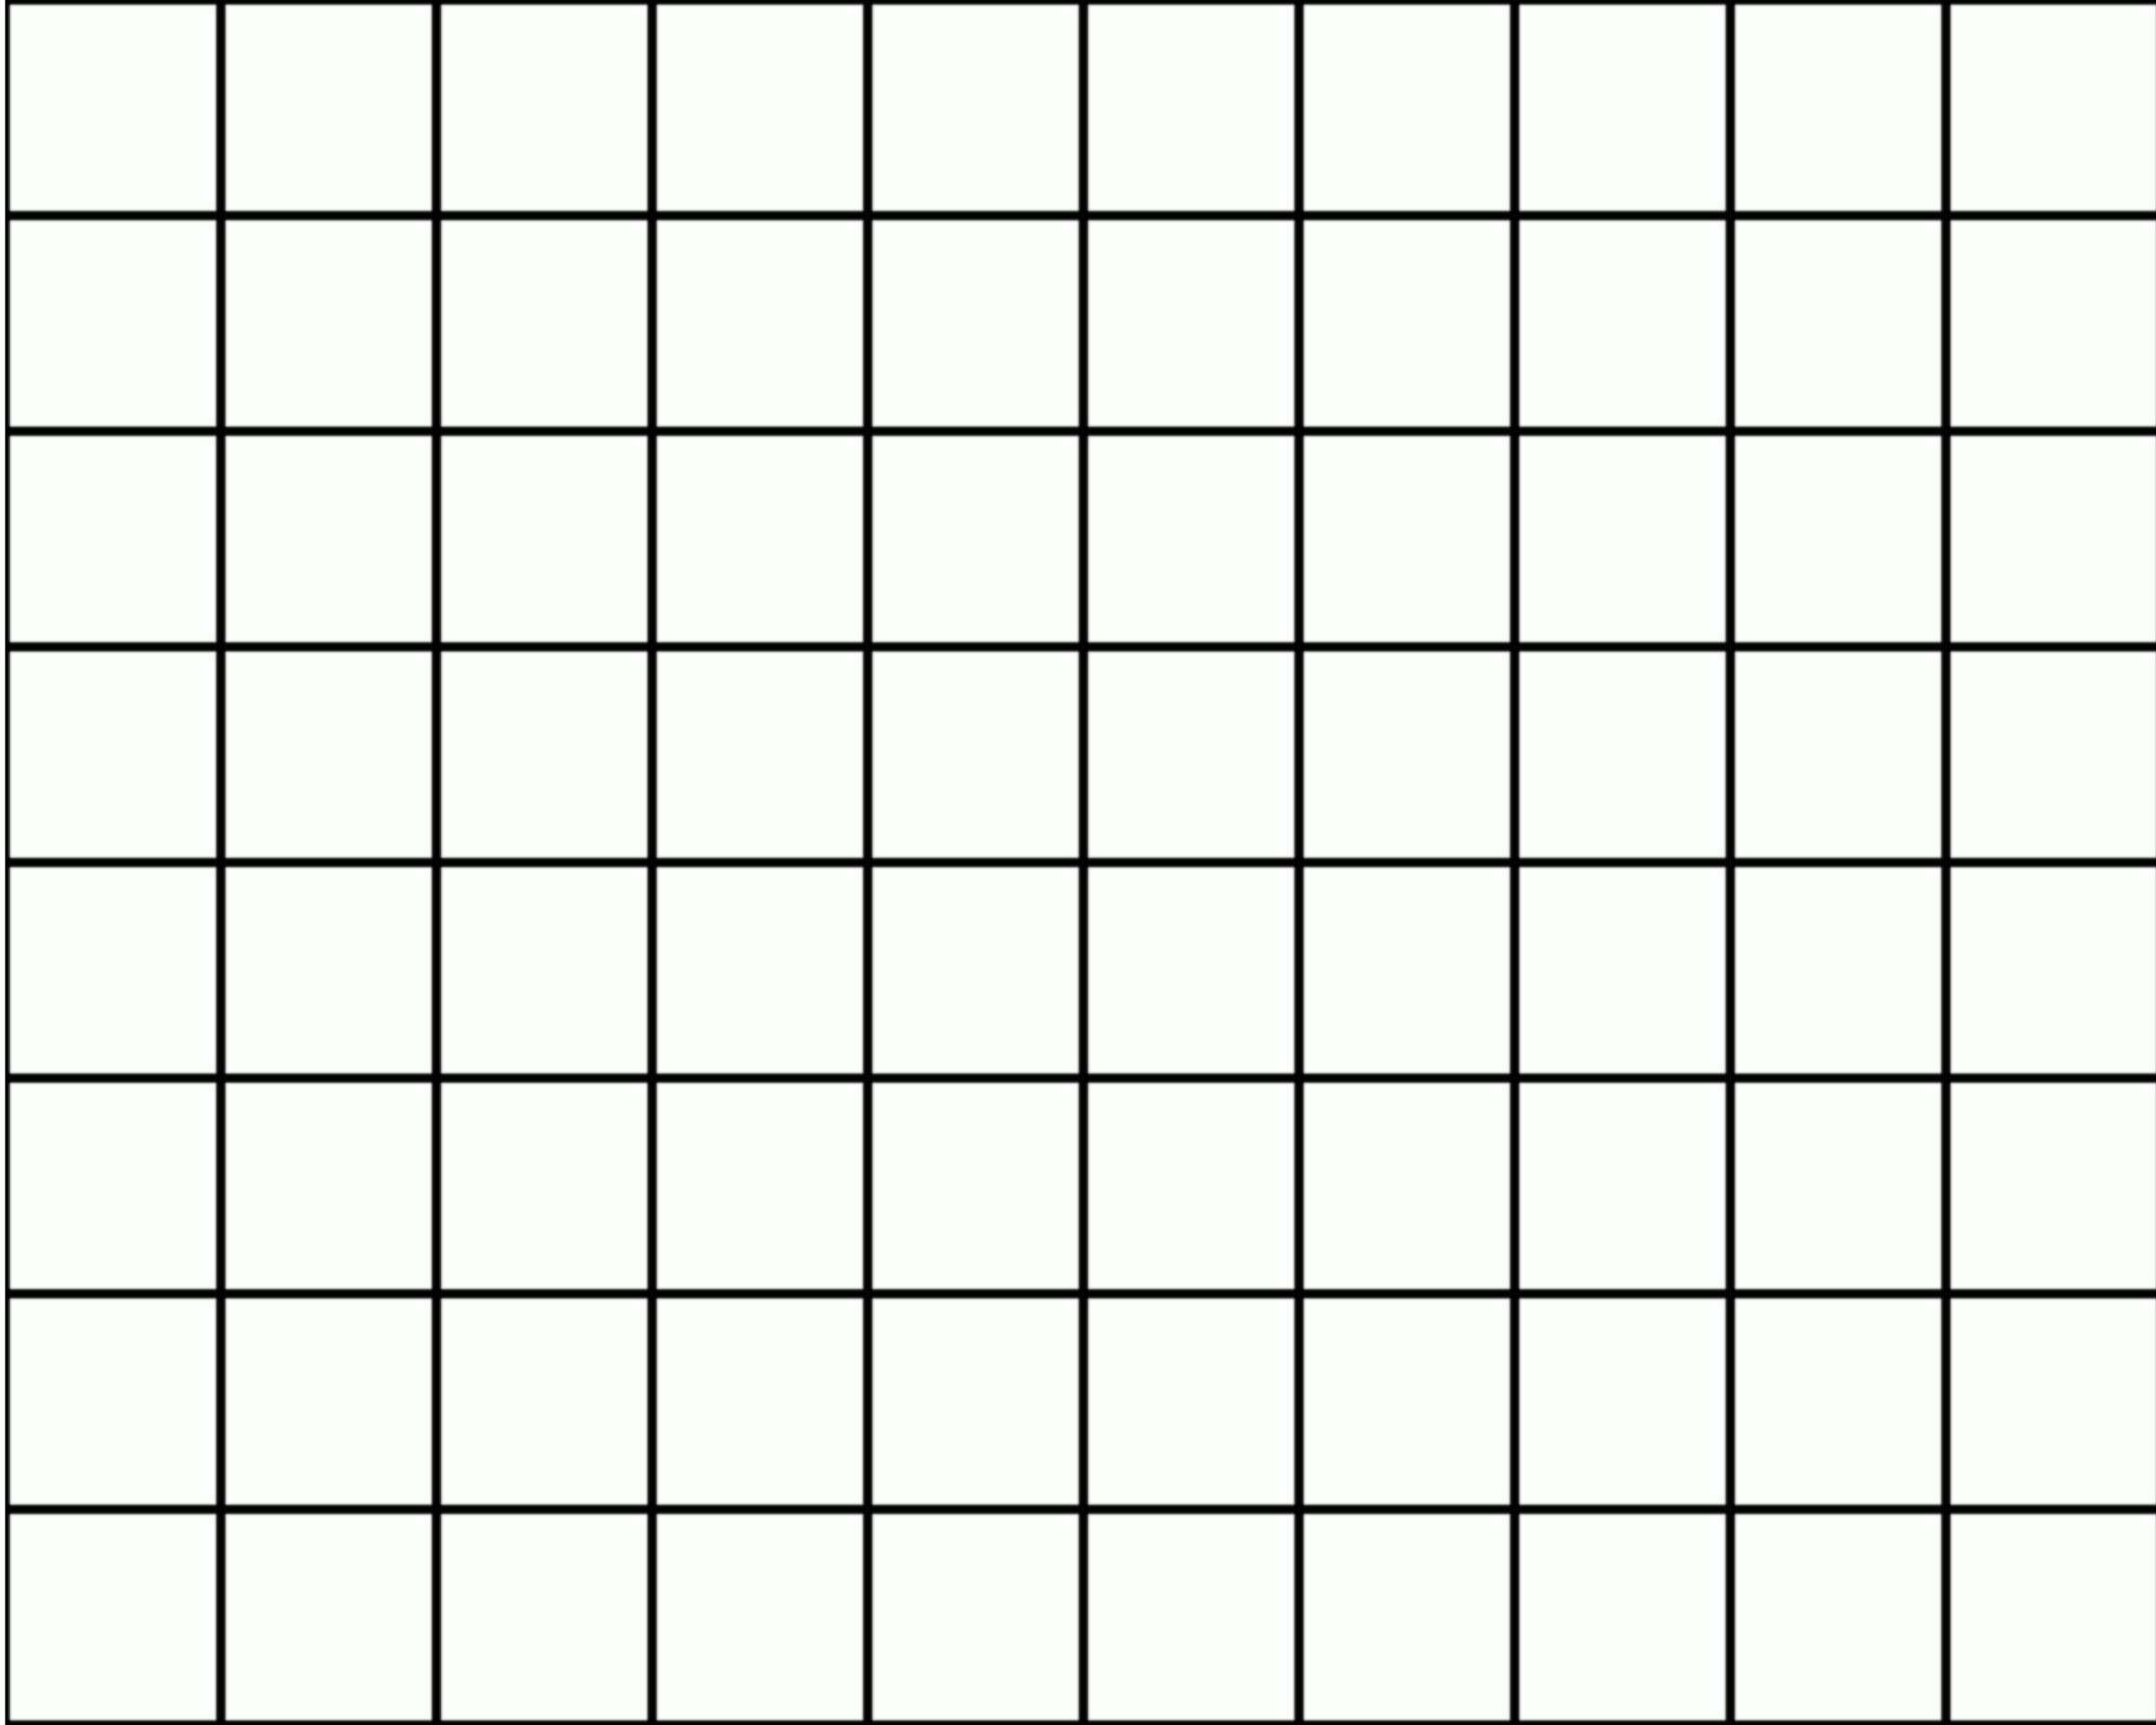
**SOLUTION?:
REDUCE ALLOCATIONS**

SOLUTION(?) MANAGED MALLOC

- » Memory pools or slabs in the runtime
- » Memcached does this
- » Multiple heaps
- » "Buddy blocks"
- » Eliminates the allocator

SOLUTION? JUST USE JEMALLOC

- » Memory contiguity is unspecified!
- » Rust: use jemalloc-specific features



SOLUTION? HARD DEFAULTS

- » Disable memory arenas
- » JEMalloc, etc
- » Disable THP

TRADEOFFS

ELIZABETH M AD 2000 THIS GREAT COUR

THANK YOU!

@NATEBERKOPEC

MORE TACTICS AT:

RAILSSPEED.COM